Build

LLMs from Scratch

What are token embeddings and why we need them?

TOKEN EMBEDDINGS

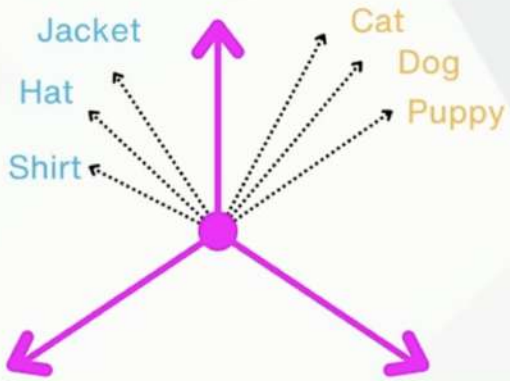

Representing Words Numerically

- Computers need numerical representation of words
- How can we represent words in numbers?

Can we assign random numbers to each word?

“cat”	→	34
“book”	→	2.9
“tablet”	→	-20
“kitten”	→	-13

The Problem With Using Random Numbers

“cat”	⟶	34
“book”	⟶	2.9
“tablet”	⟶	-20
“kitten”	⟶	-13

“cat” and “kitten” are semantically related.
However the associated numbers 34 and -13 cannot capture this relation.

What About One-Hot Encoding?

1) Create a dictionary of words
2) Assign sequential one-hot encoding to each word

"dog" → 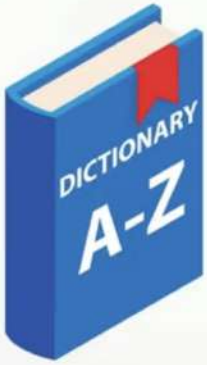

→ [0,0,0,1,0,0,0,0,0,0,0,0,0,0,.....]

The Problem With One-Hot Encoding

"dog" → 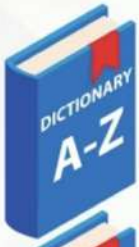

→ [0,0,0,**1**,0,0,0,0,0,0,0,0,0,0,.....]

"puppy" → 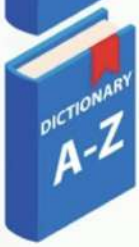

→ [0,0,0,0,0,0,0,0,0,0,0,0,0,**1**,0,.....]

One-hot encoding also fails to capture semantic relationship

Semantically Similar Words Should Have Similar Vectors

	"dog"	"cat"	"apple"	"banana"
has_a_tail	23	31	1	2
is_eatable	2	3	22	38
has_4_legs	19	21	0	0
makes_sound	12	18	0.5	0.2
is_a_pet	35	31	5	7

Vectors can capture semantic meaning

We Can Train a Neural Network To Create Vector Embedding

"dog"

"cat"

"apple"

"banana"

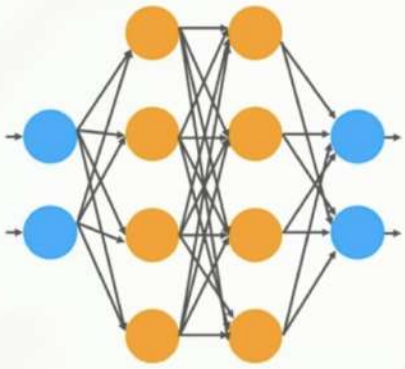

[23, 2, 19, 12, 35]

[31, 3, 21, 18, 31]

[1, 22, 0, 0.5, 5]

[2, 38, 0, 0.2, 7]